

Optimal Dub-E Scheduling

Skyler Peterson, Alex Sanchez-Stern

December 8, 2014

1 Introduction

For our final project, we decided to look to robotics for problems that SMT might be able to tackle. Since we are fortunate enough to be at a school where there is always interesting work going on across fields, we didn't have to look far. We contacted Michael Jae-Yoon Chung and Andrzej Pronobis, who are working on the Semantics Aware Robotic Assistant, more commonly known as DUB-E. DUB-E is able to traverse the CSE building, and accomplish tasks for its users, such as checking whether a particular professor is in their office. DUB-E is controlled via a web interface, from which users can request that certain tasks be accomplished, by a particular deadline.

Unfortunately, when deadlines are short and there are many tasks to accomplish, it is non-trivial to decide what task should be accomplished and when. Additionally, the expressive power of the interface to DUB-E is currently limited; Users can specify a task deadline, but cannot specify more precise timing information, such as a time in the future before which the task should not be done or multiple time periods in which a task can be accomplished. The current implementation is a simple FIFO ordering based on time of scheduling. No other temporal information is used. When a task is completed the next is begun and if the deadline has passed, an apology is sent and the next task is run.

Handling these new scheduling concerns requires a more sophisticated scheduling algorithm than the one that was previously implemented in DUB-E. We implemented this new scheduling algorithm by encoding the scheduling constraints into SMT, and then using Z3 as a backend to solve the constraints. The result is a schedule which instructs DUB-E when to tackle each task and, in some cases, how long to wait in between tasks. This encoding allows many more tasks to be completed without dropping them and also allows the robot to take advantage of time when planning task execution.

2 Overview

DUB-E's overall behavior is quite simple. It is notified of new tasks via the web interface, and its scheduler decides what tasks to do when. Each time it receives a new task, it can change how it schedules its current tasks. Tasks which are completed are removed from its schedule.

But not all tasks get completed. If DUB-E is overscheduled, it may not have enough time to complete a requested task in time. Additionally, the travel time of DUB-E is not entirely predictable. While its travel time between two locations can mostly be bounded within a range, it is always possible that it will completely fail, and take much longer to reset itself and become fully operational again.

The previous scheduling algorithm for DUB-E was quite simple. DUB-E simply maintained a FIFO queue upon which its tasks resided. New requests made via the web interface were added to the end of the queue, and when DUB-E was not currently working on a task, it pulled the next task from the front of the queue. When DUB-E does pull a task from the queue whose deadline had already passed, it notifies the user that it was unable to complete the task, and drops it from the queue. While simple and fair, this algorithm fails to provide optimal behavior in a variety of simple scenarios.

Consider the case where user one requests that DUB-E go the kitchen and check for food within the next thirty five minutes. Let's assume that it takes DUB-E fifteen minutes to complete this task, ten minutes to go to the kitchen, five to check for food. The, user two requests that DUB-E check whether Emina Torlak is in her office, within the next twenty minutes. It takes again, fifteen minutes total for DUB-E to complete the task, ten to arrive at Emina's office, and five to confirm whether or not she is there. If DUB-E addresses the tasks in a first come, first serve order, as the old scheduler would do, by the time it had finished checking for food, it would have missed the deadline for finding Emina. If instead it had reordered the tasks, checking for Emina first, it would be able to accomplish both tasks on time. This scenario can become even worse if user one gives a very large deadline and a task which takes a long time. Then if many small tasks are given which all take place in the same location but also have short deadlines, the first user has just blocked all of these task unnecessarily. In the worse case, this can allow a bad user to launch a DOS, by simply scheduling large tasks.

The DUB-E scheduling problem is not simply a classic instance of scheduling where each task takes a fixed amount of time. Consider the case where

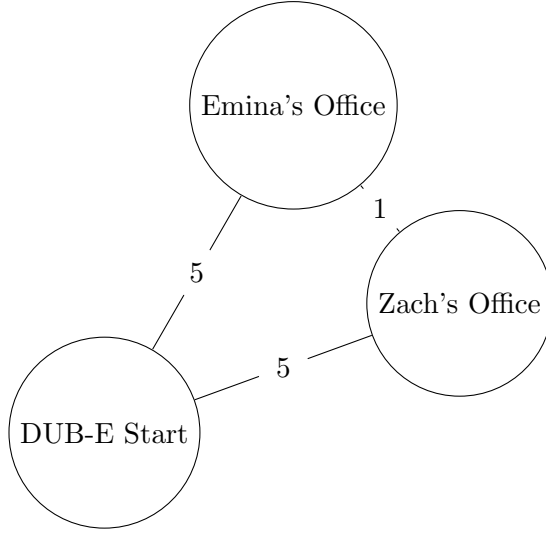


Figure 1:

DUB-E is requested to check for both Emina and Zach Tatlock in their respective offices. It might be the case that both requests should be completed in 18 minutes. If DUB-E is in another part of the building, it might take five minutes to get to both offices, and take five minutes to check them. Taken separately, each request takes ten minutes, and it is not possible to do both, since it would take twenty minutes to do both individually. But Zach's and Emina's offices are only a minute apart, both being on the fifth floor (see Figure 1), it actually is possible to complete both tasks. DUB-E can go to one office, complete the task, and travel to the other one, in only eleven minutes. Completing the second task takes a further five minutes, for a total time of sixteen minutes, well within the deadline.

It is clear from this example that a proper scheduling algorithm must consider spatial factors, as well as temporal ones. Additionally, it may be the case that not all tasks can be scheduled, and tasks have non-uniform priorities. The DUB-E team asked for multiple task priorities, that were weighted, so that sufficiently many lower priority tasks could be chosen over a higher priority task, but also include a special admin priority that would never be dropped in favor of lower tasks. An Admin task may only be dropped if conflicting with another admin task.

3 Scheduler Encoding

To tackle these unique constraints, we had to develop an encoding which could scale fairly well, while still supporting all the features the DUB-E team wanted. Since tasks can occur at a multitude of times, and include waiting times between tasks, if a task is completed and the next one belongs to an interval that has not started, a naive approach would be to ask the solver to simply give us the times at which each task is done, given the problem constraints. But solving for this many continuous variables would not scale to an even reasonable number of tasks. Instead, we discretized time into an integer counter, simplifying the issue even if we lose a little granularity, and we divided the problem space into assignment of tasks to “time steps.”

For finding the optimal schedule of N tasks, we attempt to find the best assignment of tasks to N time steps, as well as solving for a single integer for each time step indicating how long DUB-E should be waited before starting that time step. We constrain the solution to this in several ways which allow results to be a valid schedule.

Since there exists a boolean variable for each task being accomplished at each step, we add one hot-constraint so that at most one task is scheduled for each step, and each task is scheduled for at most one step. It might seem like we should require that exactly one task be accomplished at each step, and each task be accomplished at exactly one step. However, it may be the case that the tasks are over-constraining, and so not all of them can be accomplished. To account for this case, we must allow for tasks to be not scheduled for any step, and for some steps to be empty. We create a boolean variable for each time step t labeled $\text{None}@t$, which indicates whether or not **no** tasks are scheduled at that time step, and require that exactly one of the set of $\{\text{None}\} \cup \text{tasksScheduled}$ be true for each time step.

To make other constraints simpler, including the resulting schedule, we also require that all of the “None’s” be at the end of the schedule. That is, if a time step has None, then all time steps after it have None, forcing all actual tasks to the front of the time steps. (see Figure 2).

Now that we have encoded the notion of assignment of tasks to time steps into constraints, we must encode the timing constraints. But how long does traveling between tasks task? Due to the unpredictability of the robot and its environment, we don’t have a fixed bound on how long traveling between locations will take. However, we do have a notion of the longest they should take if nothing goes terribly wrong. To simplify the scheduling, we decided to ask the scheduler to only produce schedules where DUB-E will always




Figure 2:

have enough travel time if nothing goes terribly wrong. If something *does* go terribly wrong, and DUB-E falls behind schedule, we reschedule. Here, we slightly compromise on optimality, but we hope that the failure rate will be low enough that it won't be significantly worse in practice. This is one of the realities of robotics, where software and hardware must constantly be evaluating and acting upon the environment, while also being mobile. This makes almost all tasks unpredictable.

Once we have decided on an effective travel time, we create two Int variables for each time step. The first is the previously mentioned `waitBefore`, the amount of time waited after finishing the last step before starting the next one. This variable is constrained to be non-negative, but is otherwise free to be assigned whatever value makes the schedule work out best.

The second Int variable is the time at which that step is started. This variable is constrained to be exactly equal to the time of the previous time variable, plus the `waitBefore` time, plus the time taken to complete the previous task, plus the time taken to travel from the previous task to the next task. We encode the time taken to complete the previous task using the “ite” construct for each possible previous task, giving a zero value if that task was not completed last time step, and the value of the task duration if it was. We similarly encode the travel time, using every possible pair of tasks we could have traveled in between. Next, we require that if a task is done at a particular time step, we require that that time step be fully within one of the tasks intervals. That means that the task must start after the beginning of the interval, and be finished before the end.

Finally, once we have properly constrained the solver so that every solution will be a valid schedule, that the robot can execute, we attempt to find the *optimal* schedule. We decided to encode task weights as integers, where higher is better, but zero is an admin task.

First, we create a variable for each task being accomplished at some time step, and set up an implication constraint so that it can only be true if the task was completed at some time step. Then, we use a MaxSAT algorithm to determine the schedule which can fit in the most admin tasks.

Then, we run it again to find the schedule with the greatest sum weights of non-admin tasks, given that it must accomplish the admin tasks found. We reimplemented the WMax algorithm [1] because we were not able to find the tool itself. This approach is slightly flawed, since there may be multiple optimal sets of admin tasks, and one may allow more non-admin tasks, but we felt that the slight benefits of trying multiple configurations of optimal admin tasks did not outweigh the running time cost.

Once we've figured out the optimal configuration, we parse the list booleans indicating whether each task was accomplished at each time step, plus the waitBefore times of each step, into a path, and give it to the ROS interface.

4 Results

A large portion of the unknowns and design challenges in this project were related to actually trying to get the new scheduler to run in the current implementation of the SARA system which is a series of python packages built within the ROS environment. The first challenge was to meet with the SARA project leaders, Andrzej Pronobis and Michael Jae-Yoon Chung. In a few early meetings, we expressed interest in a delivery system for DUB-E and we eventually settled on a generic task scheduling problem as their current FIFO implementation was supposed to be temporary.

Several meetings had to take place after choosing a specific project while we developed designs. This included two separate problems. One was to come up with the details of the scheduling problem itself and to settle on specifically what constraints were important to the SARA team along with what was interesting for us in our own goals. The other problem involved determining how the scheduler would actually be implemented in the current SARA code base. Several problems were immediately raised. The most obvious was the package design. ROS is built as a series of network calls and is built to be almost a pure event handling system. This means we could not simply create a bunch of classes and functions and dump them in. We had to come up with service and message protocols and create independent ROS nodes that could handle reading data from the network. Figure 3 gives a basic overview of our resulting design. Another problem is that the SARA group uses a Mongo database to communicate task information. This was new to us so some additional outside knowledge was required. One problem that became apparent later in the project was the time consuming process of trying to create tests for our code. Again, ROS makes this a bit difficult as

mock databases must be created, partial SARA code must be launched, and test code gives very little feedback into problems due to the ROS protocol. Finally, using the built in visualizer was impossible in the current state of the SARA project (explained later) and the visualizer we tried to use fought us in more ways than one.

One of the main problems we faced in the implementation of the encoding was finding an appropriate MaxSAT algorithm. We initially thought that there would be a simple out-of-the-box solution with Z3. But upon looking closer, we found that the MaxSAT example in the Z3 codebase utilized special features of the C API, and was not compatible with the current syntax. Looking into the papers cited by the example, it seemed like they only addressed the problem for unweighted MaxSAT. Instead, we researched and found the vZ paper, and were able to implement some of the algorithms it described, although it took much longer than expected.

The final result was an encoder strong enough to give more flexibility to the SARA project. The new ROS package runs with 100% integration with the SARA code base in test cases. There are still a couple of things that must be done before the new scheduler can work live on the robot. As earlier discussed, the old simple was extremely simple and did not take into account any constraints except deadlines and these were only looked at during run-time. Since we added much complexity, there are a few integration issues to handle before the robot can run live. Primarily, this means providing a few more parameters in the database. We need priorities and additional time constraints which we can mock in testing and it can be simplified in live runs too. For instance, we can just always assign priorities of 0 for now. The biggest hurdle is a useful "World" description. Part of the encoder expects a World class which can give back estimates of travel time and task time. This does not exist yet, but the SARA team is currently working on this, however it is a difficult problem, due to the aforementioned chaos of real world robotics. The SARA team is expecting to fully incorporate the new SAT encoder into their project over the short time and we will continue to support them.

Finally, we would like to document one result from the project. Figure 4 shows the outputted plan of this encoder given ten tasks, chosen from three possible distinguishable tasks for each location, and ten locations, plus the start location. The red lines demonstrate the resulting path beginning at "startLoc". Table 1 shows the result in table format with tasks and wait times.

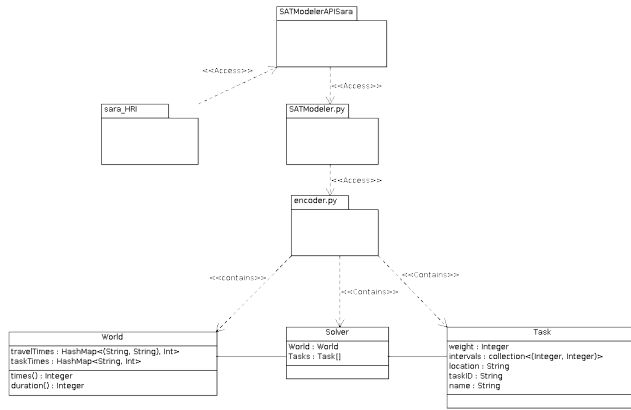


Figure 3: This is a package diagram of our ROS plugin for the SAT solver.




Figure 4: This is an example output of the sat scheduler given ten locations and ten jobs, one of three possible tasks at each location. Black lines are potential routes and red lines are the actual route. Numbers demonstrate time in seconds to travel that edge.

Event	Time to Execute (seconds)
Wait at location "startPos"	11034
Travel from location "startPos" to "Room2"	544
Due Task3 at location "Room2"	600
Wait at location "Room2"	32564
Travel from location "Room2" to "Room9"	242
Due Task1 at location "Room9"	150
Wait at location "Room9"	27021
Travel from location "Room9" to "Room3"	110
Due Task1 at location "Room3"	150
Wait at location "Room3"	27457
Travel from location "Room3" to "Room5"	713
Due Task2 at location "Room5"	400
Wait at location "Room5"	11951
Travel from location "Room5" to "Room8"	728
Due Task3 at location "Room8"	600
Wait at location "Room8"	53709
Travel from location "Room8" to "Room7"	931
Due Task2 at location "Room7"	400
Wait at location "Room7"	15881
Travel from location "Room7" to "Room1"	988
Due Task2 at location "Room1"	400
Wait at location "Room1"	46943
Travel from location "Room1" to "Room6"	752
Due Task1 at location "Room6"	150
Wait at location "Room6"	29825
Travel from location "Room6" to "Room10"	416
Due Task3 at location "Room10"	600
Wait at location "Room10"	19836
Travel from location "Room10" to "Room4"	719
Due Task2 at location "Room4"	400
Done	0

5 Project Division

For the most part we roughly divided the work into the encoding, and the interface with DUB-E. Skyler handled the interface with the robot, setting up ROS nodes and working with Michael and Andrzej to get the encoder in-

tegrated into the DUB-E codebase. This required significant domain knowledge and experience working with the "Robotic Operating System" or ROS in which the current SARA code base is implemented in. Alex handled the actual encoding of the task constraints, setting up the Z3 bindings to interface with the code, and writing modules to take in task information, encode it into a set of Z3 constraints, decode the Z3 output into a schedule, and format the schedule for DUB-E.

6 Applied Topics

This project involved mostly topics that we discussed at the beginning of the quarter, although it of course was supported by the work we did throughout the quarter. Specifically, the work on encoding different types of constraints into conjunctive normal form was paramount to the project. The encoding also made heavy use of MaxSAT, which we touched on in class, to satisfy the greatest number of tasks in cases where not all tasks could be satisfied.

References

- [1] Bjørner, Nikolaj and Phan, Anh-Dung, *vZ: Maximal Satisfaction with Z3*. <http://research.microsoft.com/en-US/people/nbjorner/scss2014.pdf> 2014.
- [2] De Moura, Leonardo and Bjørner, Nikolaj, *Z3: An Efficient SMT Solver*. Springer-Verlag, 2008.
- [3] *MongoDB*. <http://www.mongodb.org/>
- [4] *ROS: Robotic Operating System*. <http://wiki.ros.org/>
- [5] *SARA: Semantics Aware Robotic Assistant*.